

R Without Statistics

David Keyes

Contents

About the Book

This is the in-progress version of *R Without Statistics*, a forthcoming book from No Starch Press.

Since R was invented in 1993, it has become a widely used programming language for statistical analysis. From academia to the tech world and beyond, R is used for a wide range of statistical analysis.

R's ubiquity in the world of statistics leads many to assume that it is only useful to those who do complex statistical work. But as R has grown in popularity, the number of ways it can be used has grown as well. Today, R is used for:

- Data visualization
- Map making
- Sharing results through reports, slides, and websites
- Automating processes
- And much more!

The idea that R is only for statistical analysis is outdated and inaccurate. But, without a single book that demonstrates the power of R for non-statistical purposes, this perception persists.

Enter R Without Statistics.

R Without Statistics will show ways that R can be used beyond complex statistical analysis. Readers will learn about a range of uses for R, many of which they have likely never even considered.

Each chapter will, using a consistent format, cover one novel way of using R.

1. Readers will first be introduced to an R user who has done something novel and learn how using R in this way transformed their work.
2. Following this, there will be code samples that demonstrate exactly how the R user did the thing they are being profiled for.

3. Finally, there will be a summary, with lessons learned from this novel way of using R.

Written by David Keyes, Founder and CEO of R for the Rest of Us, R Without Statistics will be published by No Starch Press.

Introduction

Introduction

In early 2020, countries across the world struggled to contain the spread of COVID. One country, though, succeeded where others did not: New Zealand. There are many reasons why New Zealand was so successful in tackling COVID. One of these was R (yes, R).

How did a humble tool for data analysis help New Zealand fight COVID? R helped a team at the Ministry of Health to generate daily reports on cases throughout New Zealand. These reports (there were three each day, each for a slightly different audience) were essential in helping officials develop policies that kept New Zealand largely COVID-free. It was a big lift for a small team. Producing these reports every day with a tool like Excel would not have been feasible. As team leader Chris Knox told me, “Trying to do what we did in a point-and-click environment is not possible.” But with R, a few staff members wrote R code that they could re-run every day to produce updated reports.

The reports that the New Zealand Ministry of Health produced did not involve any complicated statistics – they were literally counts of COVID cases. The value that the team got was from everything else R can do: data analysis and visualization, report creation, and automating workflows. Many people think of R as simply a tool for statistics. But, over a quarter century since its creation, R can do much more than statistical analysis, and New Zealand used R to keep its residents safe from COVID.

I used to feel ashamed about the way I use R. As someone with an extremely non-quantitative background (I did a PhD in anthropology) who never used R in graduate school, I use R, a tool for statistical analysis, but I don’t use it for complex statistical analysis. For a long time, I felt like I wasn’t a “real” R user. Real R users, in my mind, used R for hardcore stats; I “only” used R for descriptive stats.

But eventually, I realized that, no matter what else you do in R, you have to **illuminate** your findings and **communicate** your results. And, the more you use R, the more you’ll find yourself wanting to **automate** things you used to do manually. I realize now that the things that I use R for *are* the things that everyone uses R for. R was created for statistics. But today people are just as likely to use R without statistics.

I'm excited to be your guide on this journey through the ways you can use R without statistics. If I, a qualitatively-trained anthropologist whose most complex statistical use for R is calculating averages, can find value in R, so can you. No matter your background or what you think about R right now, using R without statistics can transform your work.

Who This Book is For

This book is for you if you are either a current R user keen to explore new ways of using R or a non-R user wondering if R is right for you. I've written *R Without Statistics* so that it should make sense even if you've never written a line of R code. But if you have written many lines of R code, the book should help you learn plenty of new techniques to up your R game.

About This Book

This book shows the many ways that people use R without statistics. Each chapter focuses on one novel use of R. You'll begin by learning about R users who have transformed their work using R. You'll learn about a problem they had and how R helped them to solve it. We'll dive into their code, breaking it down to help you understand how they used R. Each chapter will conclude with a short summary, offering lessons you can take from this novel way of using R. The book has three parts:

Part 1: Illuminate

In the first part, you'll learn about ways to use R to illuminate your findings.

- **Chapter 2: Principles of Data Visualization** This chapter breaks down a visualization by Cédric Scherer and Georgios Karamanis on drought in the United States. In doing so, it shows important principles that can help you to make high-quality data visualization.
- **Chapter 3: Making Your Own Theme** This chapter shows how journalists at the BBC made a custom theme for the data visualization package known as `ggplot2`. We'll break down the `bbplot` package and in the process you'll learn how to make your own theme.
- **Chapter 4: Creating Maps** This chapter walks through the code that Abdoul Madjid used to make a map showing COVID rates in the United States in 2021. You'll learn how to use the `ggplot2` package to make high-quality maps.

- **Chapter 5: Creating High-Quality Tables** This chapter will show you how to use the `gt` package to make high-quality tables in R. Based on a conversation with Tom Mock, you'll learn to apply design principles to ensure your tables communicate effectively.

Part 2: Communicate

The second part of the book focuses on using R Markdown to communicate efficiently.

- **Chapter 6: Writing Reports in R Markdown** This chapter introduces R Markdown through a conversation with Alison Hill. A tool that allows you go from data import to final report, all in R, R Markdown can transform how you communicate. This chapter will introduce the basics to help you get started with R Markdown.
- **Chapter 7: Parameterized Reporting** One of the advantages of using R Markdown is that you can produce multiple reports at the same time using a technique called parameterized reporting. In this chapter, I speak with staff members at the Urban Institute about how they used R to produce fiscal briefs for all 50 U.S. states. In the process, you'll learn how parameterized reporting works and how you can use it.
- **Chapter 8: Making Slideshow Presentations with `xaringan`** In addition to traditional reports, R Markdown can be used to make slides. You'll come away from this chapter, which is based on my conversation with Silvia Canelón, ready to make your own presentations with the `xaringan` package.
- **Chapter 9: Building Websites with `distill`** R Markdown can also make websites. In this chapter, I speak with Matt Herman about how he used the `distill` package to make a website about COVID-19 rates in Westchester County, New York. The chapter will show you how to create your own website with R Markdown and `distill`.

Part 3: Automate

The last part of the book focuses on ways you can use R to automate your work.

- **Chapter 10: Accessing Online Data** In addition to working with data you already have, R can help you to automatically access data. This chapter shows two packages that can bring in data: `googlesheets4` for working with Google Sheets and `tidycensus` for working with United States Census Bureau data. Through conversations with Meghan Harris and Kyle Walker, you'll learn how the packages work, and how you can use them to automate the process of accessing data.

- **Chapter 11: Code Once, Run Twice: Creating Your Own Functions** One of the major benefits of R is that you can create your own functions to automate common tasks. In this chapter, I show a few example functions that I and others have made. You'll come away ready to make your own R functions.
- **Chapter 12: Bundle Your Functions Together in Your Own R Package** Once you have a set of functions that you use regularly, you'll want to bundle them into a package. Doing so makes it easy for you and others to use the code you've written. I speak with Travis Gerke and Garrick Aden-Buie about how they created packages to improve the work of researchers at the Moffitt Cancer Center. This chapter will set you up to make your own R package.

Before we dive into the book, I have a favor to ask. This book is called *R Without Statistics*. But it's not meant to be taken literally. Of course it's true that if you're making a graph you're using statistics. Before you start typing an angry email, please know that R Without Statistics is a mindset, not a statement meant to be taken literally. We're all using R with statistics already. Let's learn to use R without statistics.

An R Programming Crash Course

R has a well-earned reputation for being hard to learn. Especially for those who come to R without programming experience, it can be hard to figure out how things work. This chapter is designed to help those who have never used R before. I'll start from scratch, showing you what you need to download in order to use R, and how to work with data using functions, objects, packages, and projects. If you have some experience with R, feel free to skip this chapter. But if you're just starting out, this chapter will help you understand the basics, and help you make sense of the rest of the book.

Getting Set Up

One of the more confusing things for people just starting out is that you need two pieces of software in order to use R. The first is R itself, which provides the underlying computational tools that make R work. The second is RStudio, which makes working with R much easier. The best way to understand the relationship between R and RStudio is with this analogy from the book *Modern Dive* by Chester Ismay and Albert Kim. R is the engine that makes your work with data go. RStudio is like a dashboard that makes it easier to work with your data by providing a more user-friendly interface.

Let's download each piece and get started. To download R, go to <https://cloud.r-project.org/> and choose your operating system, as seen in Figure ??

Once you download and install R, open it and you can work at the command line. For example, I can type $2 + 2$, hit enter, and I will see 4, as seen in Figure ??.

Simple math problems are only the start; you can do pretty much anything in R. No matter what you're planning to do, you're probably not super impressed with the R interface. A few brave souls work only using the command line we're looking at, but most do not. RStudio is where most R coders do their

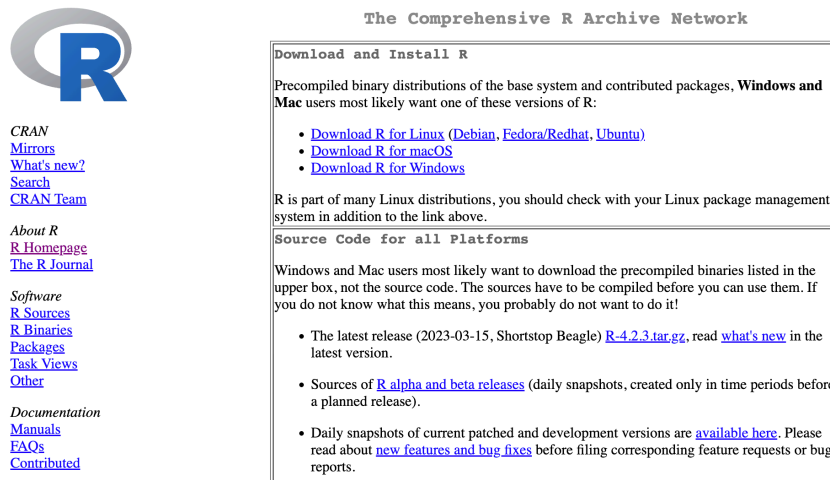


Figure 1: The Comprehensive R Archive Network where you can download R

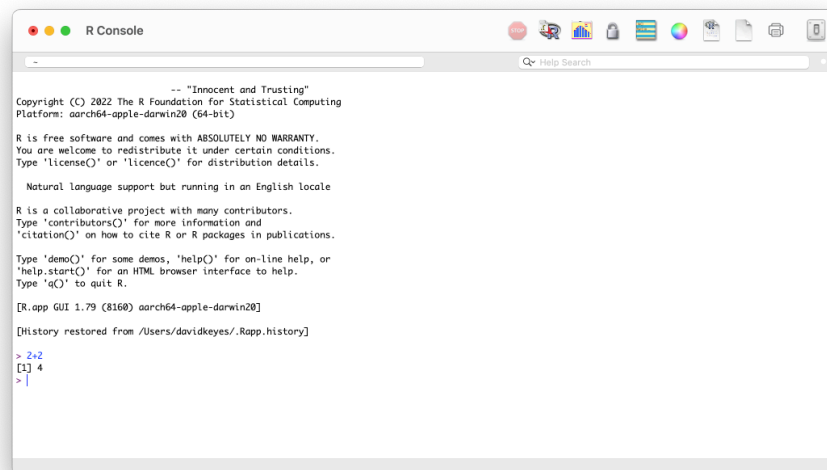


Figure 2: The R console

work. RStudio is like a skin that lives on top of R itself. It doesn't provide new functionality to R, but it wraps R in a much more user-friendly interface, providing a way to see your files, outputs, and more. You can download RStudio at <https://posit.co/download/rstudio-desktop/>. Install RStudio as you would any other app and open it up. RStudio has several panels. The first time you open RStudio, you'll see these the three panels shown in Figure ?? below.

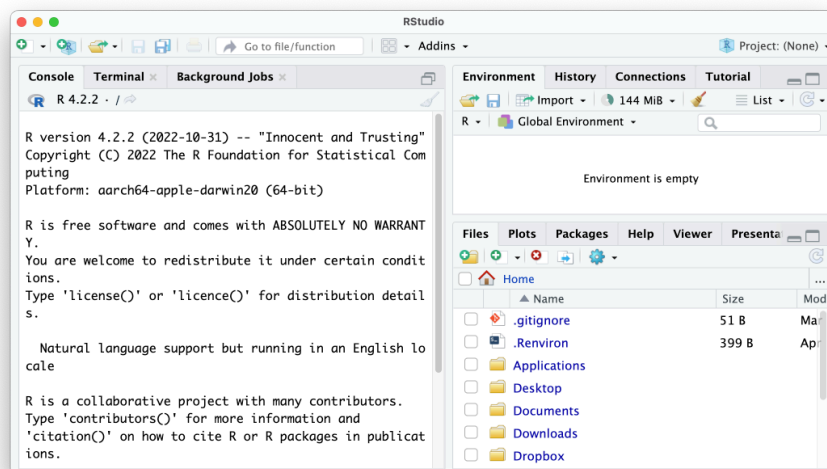


Figure 3: The RStudio editor

The left side panel should look familiar. It is what we saw when working in R. This is known as the **console**. You'll use it to add code and see results. This panel, like the others we'll discuss, has several tabs (terminal and background jobs) for more advanced usages. For now, we'll stick to the default tab. Let's look at the bottom right panel next. This **files** panel shows all of the files on my computer. Finally, the top right panel shows my **environment**. The environment shows the objects (discussed below) I have available to me when working in RStudio. There's one more panel that you'll typically have when working in RStudio. But to make it appear, we need to create an R script file.

R Script Files

If you work in the console, either in RStudio or in R itself, you don't have a record of your code. Say you sit down today and write code to import your data, analyze it, and make some graphs. You don't want to have to recreate that code from scratch tomorrow. The way to save your code is by using files.

Files allow you to save all of the code you have written. There are two types of files we'll discuss in this book:

1. R script files, which only contain code.
2. R Markdown files, which contain code combined with text.

We'll talk about R Markdown files starting in Chapter ???. Let's start with R script files, which use the `.R` extension. To create an R script file, go to `File > New File > R Script`. When you create a new R script file, you'll now have a fourth panel in the top left, which you can see in Figure ???. I'll save this file in my `Documents` folder as `sample-code.R`.

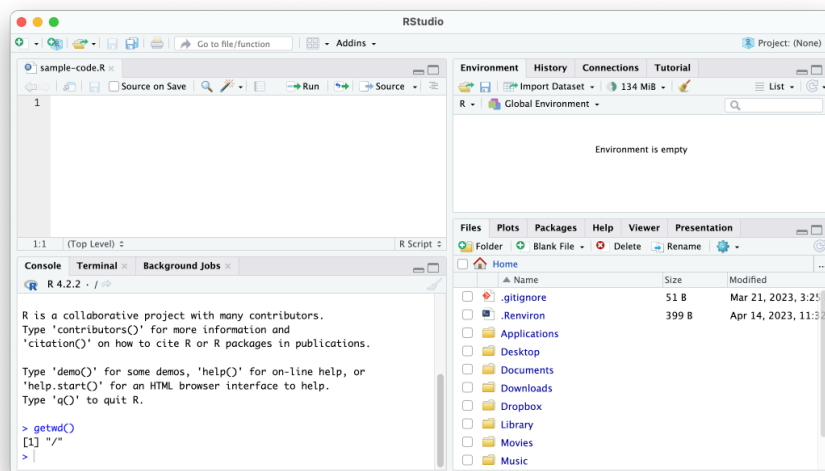


Figure 4: RStudio with four panels

I can now use the same syntax in my R script file that I did when working in just R. If I type `2 + 2` in the R script file and hit the **Run** button, 4 will show up in the console pane. If you're looking to learn R, it's probably not to help you figure out the answer to `2 + 2`. Instead, you probably want to read in your own data and do analysis on it. Let's work with some real data.

Working with Data

To explain how you work with data in R, we need go on a bit of a detour. We'll make stops to discuss RStudio functions, objects, packages, and projects before we import and take a look at our data.

Conceptually, working with data in R is very different than working with data in a tool like Excel. In Excel, your data and any analysis you do on it all live in the same place: a spreadsheet. With R, you typically have data that lives in some external source (for example, an Excel spreadsheet or a CSV file). In order to work with this data in R, you have to run code to import it. It's only once you've run this code, which is made up of functions, that you have the data available in R.

Functions

Let's say I have a CSV file called `population-by-state.csv` in my `Documents` folder that I want to import to R. To import it into R, you might think to add a line like this in the `sample-code.R` file:

```
read.csv(file = "/Users/davidkeyes/Documents/population-by-state.csv")
```

This line shows the `read.csv()` function. Functions in R are pieces of code that you can run to do specific things. Functions have a name and arguments, which are surrounded by parentheses. Looking at the `read.csv()` function, the name, which appears before the open parentheses, is `read.csv`. Within the parentheses, we have the text `file = "Documents/population-by-state.csv"`. Here we can see the argument `file`. The text after the equals sign gives the location of the file we want to read in. Arguments work in this way: the argument name, followed by the equals sign, followed by some value. This allows us to do something general (like importing a CSV) while allowing us to choose the specific file to run the function on. Functions can have multiple arguments as well, each of which is separated by a comma. For example, this would read in the same file, but skip the first row.

```
read.csv(file = "/Users/davidkeyes/Documents/population-by-state.csv",
         skip = 1)
```

At this point, you might think to run the code in order to import your data. You can do so by selecting the line of code and hitting the Run button (or using the keyboard shortcut Command/Control + Enter on Mac/Windows). Running this code causes this text show up in the console pane.

```
#>      rank      State      Pop  Growth  Pop2018
#> 1      1      California 39613493  0.0038 39461588
#> 2      2          Texas 29730311  0.0385 28628666
#> 3      3      Florida 21944577  0.0330 21244317
#> 4      4      New York 19299981 -0.0118 19530351
#> 5      5  Pennsylvania 12804123  0.0003 12800922
```

```

#> 6      6      Illinois 12569321 -0.0121 12723071
#> 7      7      Ohio 11714618  0.0033 11676341
#> 8      8      Georgia 10830007  0.0303 10511131
#> 9      9      North Carolina 10701022  0.0308 10381615
#> 10     10     Michigan 9992427  0.0008 9984072
#> 11     11     New Jersey 8874520 -0.0013 8886025
#> 12     12     Virginia 8603985  0.0121 8501286
#> 13     13     Washington 7796941  0.0363 7523869
#> 14     14     Arizona 7520103  0.0506 7158024
#> 15     15     Tennessee 6944260  0.0255 6771631
#> 16     16     Massachusetts 6912239  0.0043 6882635
#> 17     17     Indiana 6805663  0.0165 6695497
#> 18     18     Missouri 6169038  0.0077 6121623
#> 19     19     Maryland 6065436  0.0049 6035802
#> 20     20     Colorado 5893634  0.0356 5691287
#> 21     21     Wisconsin 5852490  0.0078 5807406
#> 22     22     Minnesota 5706398  0.0179 5606249
#> 23     23     South Carolina 5277830  0.0381 5084156
#> 24     24     Alabama 4934193  0.0095 4887681
#> 25     25     Louisiana 4627002 -0.0070 4659690
#> 26     26     Kentucky 4480713  0.0044 4461153
#> 27     27     Oregon 4289439  0.0257 4181886
#> 28     28     Oklahoma 3990443  0.0127 3940235
#> 29     29     Connecticut 3552821 -0.0052 3571520
#> 30     30     Utah 3310774  0.0499 3153550
#> 31     31     Puerto Rico 3194374  0.0003 3193354
#> 32     32     Nevada 3185786  0.0523 3027341
#> 33     33     Iowa 3167974  0.0061 3148618
#> 34     34     Arkansas 3033946  0.0080 3009733
#> 35     35     Mississippi 2966407 -0.0049 2981020
#> 36     36     Kansas 2917224  0.0020 2911359
#> 37     37     New Mexico 2105005  0.0059 2092741
#> 38     38     Nebraska 1951996  0.0137 1925614
#> 39     39     Idaho 1860123  0.0626 1750536
#> 40     40     West Virginia 1767859 -0.0202 1804291
#> 41     41     Hawaii 1406430 -0.0100 1420593
#> 42     42     New Hampshire 1372203  0.0138 1353465
#> 43     43     Maine 1354522  0.0115 1339057
#> 44     44     Montana 1085004  0.0229 1060665
#> 45     45     Rhode Island 1061509  0.0030 1058287
#> 46     46     Delaware 990334  0.0257 965479
#> 47     47     South Dakota 896581  0.0204 878698
#> 48     48     North Dakota 770026  0.0158 758080
#> 49     49     Alaska 724357 -0.0147 735139
#> 50     50 District of Columbia 714153  0.0180 701547
#> 51     51     Vermont 623251 -0.0018 624358

```

```

#> 52      52      Wyoming      581075      0.0060      577601
#>      Pop2010 growthSince2010 Percent      density
#> 1  37319502      0.0615  0.1184      254.2929
#> 2  25241971      0.1778  0.0889      113.8081
#> 3  18845537      0.1644  0.0656      409.2229
#> 4  19399878     -0.0051  0.0577      409.5400
#> 5  12711160      0.0073  0.0383      286.1704
#> 6  12840503     -0.0211  0.0376      226.3967
#> 7  11539336      0.0152  0.0350      286.6944
#> 8   9711881      0.1151  0.0324      188.3054
#> 9   9574323      0.1177  0.0320      220.1041
#> 10  9877510      0.0116  0.0299      176.7351
#> 11  8799446      0.0085  0.0265     1206.7609
#> 12  8023699      0.0723  0.0257      217.8776
#> 13  6742830      0.1563  0.0233      117.3249
#> 14  6407172      0.1737  0.0225       66.2016
#> 15  6355311      0.0927  0.0208      168.4069
#> 16  6566307      0.0527  0.0207      886.1845
#> 17  6490432      0.0486  0.0203      189.9644
#> 18  5995974      0.0289  0.0184       89.7419
#> 19  5788645      0.0478  0.0181      624.8518
#> 20  5047349      0.1677  0.0176       56.8653
#> 21  5690475      0.0285  0.0175      108.0633
#> 22  5310828      0.0745  0.0171       71.6641
#> 23  4635649      0.1385  0.0158      175.5707
#> 24  4785437      0.0311  0.0147       97.4271
#> 25  4544532      0.0181  0.0138      107.0966
#> 26  4348181      0.0305  0.0134      113.4760
#> 27  3837491      0.1178  0.0128       44.6872
#> 28  3759944      0.0613  0.0119       58.1740
#> 29  3579114     -0.0073  0.0106      733.7507
#> 30  2775332      0.1929  0.0099       40.2918
#> 31  3721525     -0.1416  0.0095      923.4964
#> 32  2702405      0.1789  0.0095       29.0195
#> 33  3050745      0.0384  0.0095       56.7158
#> 34  2921964      0.0383  0.0091       58.3059
#> 35  2970548     -0.0014  0.0089       63.2186
#> 36  2858190      0.0207  0.0087       35.6808
#> 37  2064552      0.0196  0.0063       17.3540
#> 38  1829542      0.0669  0.0058       25.4087
#> 39  1570746      0.1842  0.0056       22.5079
#> 40  1854239     -0.0466  0.0053       73.5443
#> 41  1363963      0.0311  0.0042      218.9678
#> 42  1316762      0.0421  0.0041      153.2674
#> 43  1327629      0.0203  0.0040       43.9167
#> 44   990697      0.0952  0.0032        7.4547

```

```
#> 45 1053959      0.0072 0.0032 1026.6044
#> 46  899593      0.1009 0.0030  508.1242
#> 47  816166      0.0985 0.0027   11.8265
#> 48  674715      0.1413 0.0023   11.1596
#> 49  713910      0.0146 0.0022    1.2694
#> 50  605226      0.1800 0.0021 11707.4262
#> 51  625879     -0.0042 0.0019   67.6197
#> 52  564487      0.0294 0.0017    5.9847
```

This is R confirming that it read in the CSV file and showing us the data within it. You might think you are ready to work with your data in R. But in fact all you've done at this point is **display** the result of running the code that imports your data. To use the data again, you need to **save** the result of running the code to an object.

Objects

To save your data for reuse, you need to create an object. To do so, you would add to your data importing syntax from above.

```
population_data <- read.csv(file = "/Users/davidkeyes/Documents/population-by-state.csv")
```

The second half of this code is what we used above, but we've added to it. In the middle you will see this: `<-`. Known as the assignment operator, it takes what follows it and assigns it to the item on the left. To the left of the assignment operator is `population_data`. This is an **object**. Put together, the whole line reads in the CSV and assigns it to an object called `population_data`. If you run this line of code, you will now see `population_data` in your environment pane, as in Figure ??.

This is confirmation that your data import worked and you have the `population_data` object ready for future use. Now, instead of having to rerun the code to import the data, I can simply type `population_data`, run that line, and I'll see the same output as above. Data imported to an object is known as a **data frame**.

Packages

The `read.csv()` function that we've used up to this point is one of a set of functions that come from what is known as base R. They are built into R and you simply have to type the name of the function to use it. However, one of the benefits of R being an open source language is that anyone create their own

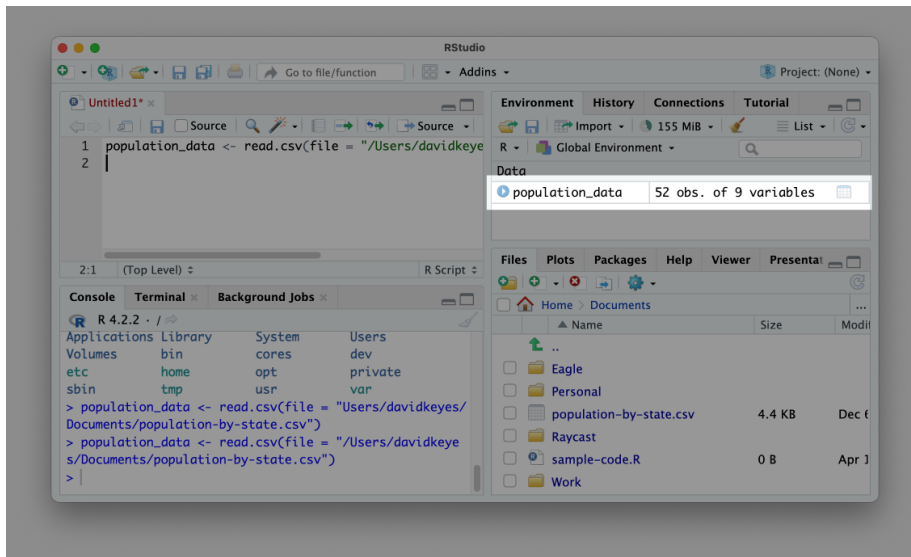


Figure 5: An object in our environment pane

code and share it with others. R users around the world make what are called **packages**, which provide code to do specific things.

The best analogy for understanding packages also comes from the *Modern Dive* book. The functionality in base R is like that built into a phone. A phone can do a lot on its own. But you usually want to install apps on your phone to do specific things. Packages are like apps, giving you specific functionality that doesn't come built into base R.

You can install packages using the `install.packages()` function. For example, to install the `tidyverse` package, which provides a range of functions for data import, cleaning, analysis, visualization, and more, you would type `install.packages("tidyverse")`. I typically enter this code in the console because you only need to install a package once on your computer and so I know I won't need to rerun this code.

To confirm that the `tidyverse` package has been installed correctly, click on the packages tab on the bottom right panel. Search for `tidyverse` and you should see it pop up, as in Figure ??.

Now that we've installed the `tidyverse` package, let's use it. While you only need to install packages once per computer, you need to load packages each time you restart RStudio. You can only use functions from the `tidyverse` package if you first run the line `library(tidyverse)`. I'll go back to my `sample-code.R` file and re-import my data using a function from the `tidyverse` package.

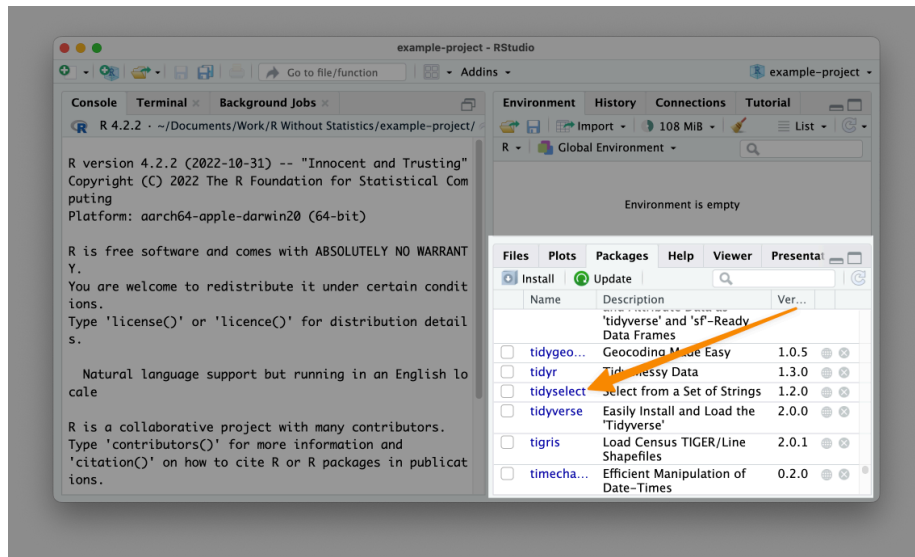


Figure 6: Confirmation that the tidyverse package is installed on my computer

```
library(tidyverse)
```

```
population_data_2 <- read_csv(file = "/Users/davidkeyes/Documents/population-by-state.csv")
```

At the top of my script I load the `tidyverse`. Then, I use the `read_csv()` function (note the `_` in place of the `.`) to import my data. This alternate function to import CSV files achieves the same goal of creating an object called `population_data_2`. If we type `population_data_2` and run the code (either by using the run button or the keyboard shortcut) you will see the output in the console.

```
#> # A tibble: 52 x 9
#>   rank State      Pop Growth Pop2018 Pop2010
#>   <dbl> <chr>    <dbl>   <dbl>   <dbl>   <dbl>
#> 1     1 California 39613493 0.0038 39461588 37319502
#> 2     2 Texas    29730311 0.0385 28628666 25241971
#> 3     3 Florida 219445777 0.033 21244317 18845537
#> 4     4 New York 19299981 -0.0118 19530351 19399878
#> 5     5 Pennsylvania 12804123 0.0003 12800922 12711160
#> 6     6 Illinois 12569321 -0.0121 12723071 12840503
#> 7     7 Ohio    11714618 0.0033 11676341 11539336
#> 8     8 Georgia 10830007 0.0303 10511131 9711881
#> 9     9 North Carolina 10701022 0.0308 10381615 9574323
#> 10    10 Michigan 9992427 0.0008 9984072 9877510
```

```
#> # i 42 more rows
#> # i 3 more variables: growthSince2010 <dbl>, Percent <dbl>,
#> #   density <dbl>
```

What we see is slightly different from what we saw above using the `read.csv()` function. R describes the output as a **tibble** and only shows us the first 10 rows. This slightly different output occurs because `read_csv()` imports the data not as a data frame, but as a tibble. Both are used to describe rectangular data like what you would see in a spreadsheet. While there are some small differences between data frames and tibbles, I'll use the terms interchangeably in this book.

RStudio Projects

So far, we've imported a CSV file from the Documents folder. But the path to the file on my computer was `/Users/davidkeyes/Documents/population-by-state.csv`. Since others will not have this exact location on their computer, if they try to run my code, it won't work. There's a solution to this problem, and it's called **RStudio projects**.

By working in a project, you can use what are known as **relative paths** to your files. Instead of having to write out `read_csv(file = "/Users/davidkeyes/Documents/population-by-state.csv")`, you can put the CSV file in your project and then call it using `read_csv(file = "population-by-state.csv")`. This makes it easier for you, and enables others to use your code.

To create a new RStudio project, go to `File > New Project`. Select either `New Directory` or `Existing Directory` and choose where to put your project. If you choose `New Directory`, you'll need to specify that you want to create a new project. I'll do this and then choose a name for the new directory and where it should live. As seen in Figure ??, you can leave the two checkboxes that ask about creating a git repository and using `renv` unchecked (these are for more advanced purposes).

Having now created this project, there are two major differences in RStudio's appearance:

First, the files pane no longer shows every file on my computer, but instead only shows files in the `example-project` directory. Right now that's just the `example-project.Rproj` file that indicates the folder contains a project. Second, at the very top right of RStudio, you can see the name of the `example-project` project (it had previously said `Project: (None)`). If you want to make sure you're working in a project, make sure you see its name here. Both of these changes can be seen in Figure ?? below.

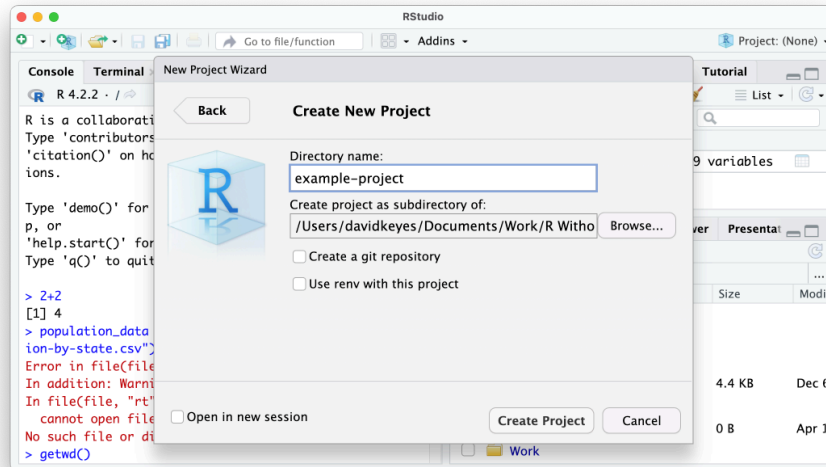


Figure 7: The RStudio prompt to create a new project

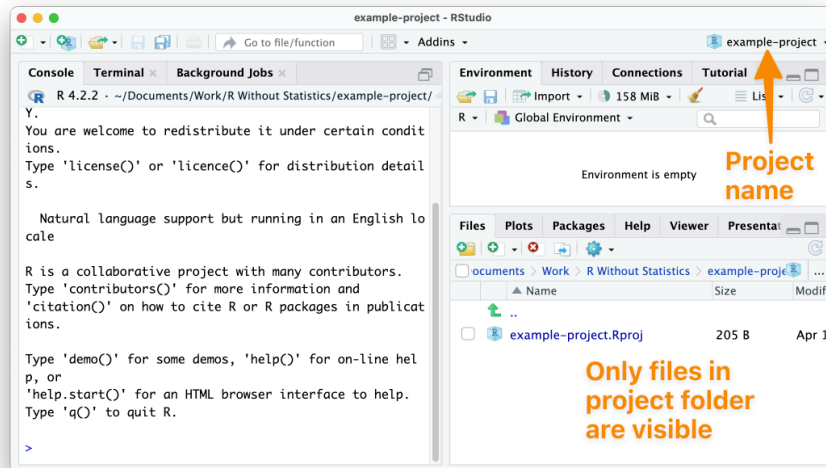


Figure 8: RStudio with an active project

Now that I've created a project, I'll use the Finder on my Mac computer to copy the `population-by-state.csv` file into the `example-project` directory. Once I do this, I can see it in the RStudio files pane, as in Figure ??.

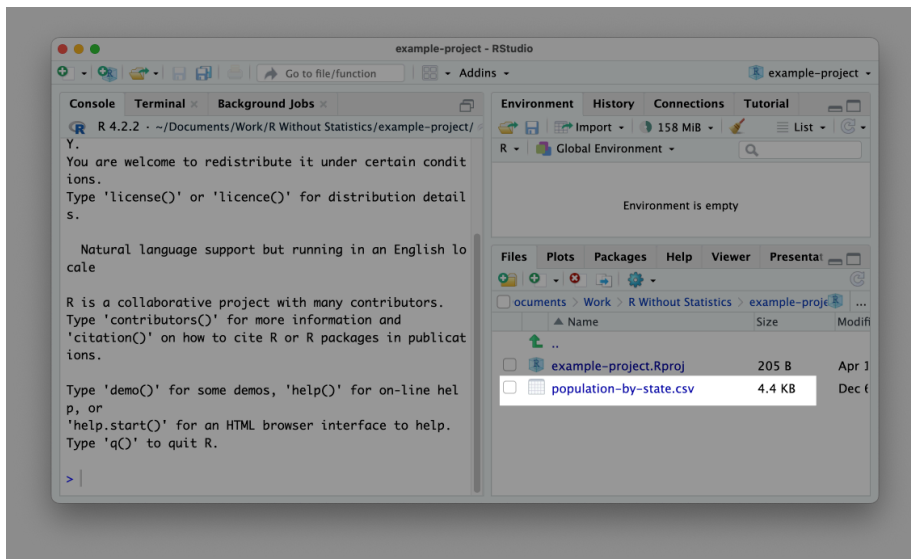


Figure 9: A CSV file visible in the files pane in RStudio

With this CSV file in my project, I can now import it more easily. As before, I'll start by loading the `tidyverse` package. After that, I can remove the reference to the `Documents` folder and import my data by simply using the name of the file:

```
library(tidyverse)

population_data_2 <- read_csv(file = "population-by-state.csv")
```

I'm able to import the `population-by-state.csv` file in this way because the RStudio project sets the **working directory** to be the root of my project. With the working directory set in this way, all references to files are relative to the `.Rproj` file at the root of the project (this is where the name relative paths comes from). Now that we're working in a project, anyone can run this code because it imports the data from a location that they are guaranteed to have on their computer.

Data Analysis with R

Now that we've imported data, let's do a bit of analysis on it. Below is a code snippet that calculates the mean population of all states using the `summarize()` function.

```
summarize(.data = population_data_2,  
          mean_population = mean(Pop))
```

You can see that I'm using `population_data_2` with the `.data` argument, telling the `summarize()` function (which comes from the `tidyverse`) to use that data frame. The second half of the code creates a new variable called `mean_population`, which is calculated by using the `mean()` function on the `Pop` variable. Running this code will return a tibble with a single variable (`mean_population`) that is of type double (meaning numeric) and has a value of 6433422, the mean population of all states.

```
#> # A tibble: 1 x 1  
#>   mean_population  
#>   <dbl>  
#> 1      6433422.
```

This is a basic example of data analysis, but you can do a lot more with the `tidyverse`. One advantage of working with the `tidyverse` is that it uses what's known as the **pipe** for multi-step operations. The `tidyverse` pipe, which is written with the text `%>%`, allows us to break steps into multiple lines (the functionally equivalent so-called native pipe uses the text `|>`). For example, I can rewrite the code above to do the same thing using the pipe.

```
population_data_2 %>%  
  summarize(mean_population = mean(Pop))
```

This code says: start with the `population_data_2` data frame, then run the `summarize()` function on it, creating a variable called `mean_population` by calculating the mean of the `Pop` variable.

The pipe becomes even more useful when we use multiple steps. Let's say, for example, we want to calculate the mean population of the five largest states. The code below adds a line that uses the `filter()` function (also from the `tidyverse`) to only include states where the `rank` variable (which is the rank of the total population size of all states) is less than or equal to five (in other words, rank one through five). Then, it uses `summarize()` function as we did before.

```
population_data_2 %>%
  filter(rank <= 5) %>%
  summarize(mean_population = mean(Pop))
```

Running this code shows us the mean population of the five largest states.

```
#> # A tibble: 1 x 1
#>   mean_population
#>   <dbl>
#> 1      24678497
```

Combining functions using the pipe lets us do multiple things on our data in a way that keeps our code readable and easy to understand. We've introduced only two functions for analysis at this point, but the **tidyverse** has many functions that enable you to do nearly anything you could hope to do with your data. In fact, while I've been referring to the **tidyverse** as a single package, it is actually a collection of packages that do data importing, analysis, visualization, and more. The book *R for Data Science* by Hadley Wickham, Mine Çetinkaya-Rundel, and Garrett Grolemund is the bible of **tidyverse** programming and worth reading for more details on how its many packages work. In this book, I'll introduce you to a number of packages, but because of how useful it is, the **tidyverse** will appear in every single piece of R code I write.

How to Get Help

Now that you've learned about the basics of how R works, you're probably ready to dive in. When you do, you're going to encounter errors. Everyone does, and it's just part of working in R. Learning how to get help when you do run into issues is a key part of learning to use R successfully. There are two main strategies you can use to get unstuck.

The first is to read the documentation for functions. To access the documentation for any function, simply type the `?` and then the name of the function. For example, if I run the line `?read.csv`, I will see the documentation pop up in the bottom right panel, as in Figure ??.

Help files can be a bit hard to decipher but at their core, they tell you what package the function comes from, what it does, its arguments, and some examples of how to use it. For additional guidance on reading documentation, I recommend the appendix of Kieran Healy's book *Data Visualization: A practical introduction* (a free online version is available at <https://socviz.co/appendix.html>).

In addition to providing help files in RStudio, many R packages have documentation websites. I find these easier to read and tend to use them when I am

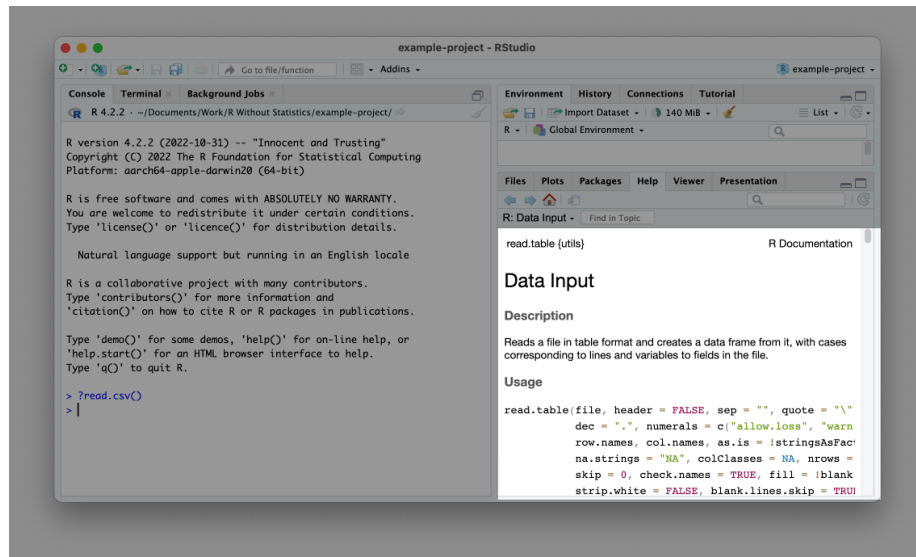


Figure 10: The documentation for the ‘read.csv()’ function

confused about how to use a function. In addition, many packages have longer articles known as vignettes that provide an overview of how the package works. Reading these can help you see how individual functions can be used in the context of a larger project. Every package I discuss in this book has a good documentation website.

In Conclusion: Invest Time Now to Learn R and Save Time Later

Getting started with R can be challenging. I myself experienced many frustrations early on that I’ve since realized are quite common. This chapter has, hopefully, helped you to see how you can get started with R. Understanding how functions, objects, packages, and projects work is key to ensuring that you can successfully use R to work with your data.

If R feels challenging, just know that it will get better with time. Best of all, the time you invest in learning to use R will repay itself many times over. My favorite example to show this is one I discuss in Chapter ??, which discusses a technique called parameterized reporting to automatically produce dozens, hundreds, or even thousands of reports at once. At my company, R for the Rest of Us, we worked with a client to produce reports on demographics and housing data for each of the 170 plus towns and counties in the state of Connecticut. Doing this by hand would have taken the client hundreds of hours. Using R,

we were able to automate the process so they can generate the reports simply by running code. If you make multiple reports by hand, think of the hours you're spending. Reframe the time it takes to learn R as an investment in never having to do this manual labor again. When you're struggling to make sense of an inscrutable error message, you will get frustrated. But, with all of the time you'll save, I promise that it is worth the effort to learn R.

Illuminate

Principles of Data Visualization

In the spring of 2021, nearly all of the American West was in a drought. By April of that year, officials in Southern California had declared a water emergency, citing unprecedented conditions.

This wouldn't have come as news to those living in California and other Western states. Drought conditions like those in the West in 2021 are becoming increasingly common. Yet communicating the extent of problem remains difficult. How can we show the data in a way that accurately represents it while making it compelling enough to get people to take notice? This was the challenge that data-visualization designers Cédric Scherer and Georgios Karamanis took on in the fall of 2021. Working with the magazine *Scientific American* to create a data visualization of drought conditions over the last two decades in the United States, they turned to the ggplot2 package to transform what could have been dry data (pardon the pun) into a visually arresting and impactful graph.

In this chapter, I show how Scherer and Karamanis made their data visualization. We begin by looking at why the data visualization is effective. Next, we talk about the grammar of graphics, a theory to make sense of graphs that underlies the ggplot2 package that Scherer, Karamanis, and millions of others use to make data visualization. We then return to the drought graph, recreating it step-by-step using ggplot2. In the process, we pull out some key principles of high-quality data visualization that you can use to improve your own work.

The Drought Visualization

There was nothing unique about the data that Scherer and Karamanis used. Other news organizations had relied on the same data, from the National Drought Center, in their stories. But Scherer and Karamanis visualized it in a way that it both grabs attention and communicates the scale of the phenomenon. Figure ?? shows a section of the final visualization. Showing four

regions over the last two decades, the increase in drought conditions, especially in California and the Southwest, is made apparent.

To understand why this visualization is effective, let's break it down into pieces. At the broadest level, the data visualization is notable for its minimalist aesthetic. There are, for example, no grid lines and few text labels, as well as little text along the axes. What Scherer and Karamanis have done is remove what statistician Edward Tufte, in his 1983 book *The Visual Display of Quantitative Information*, calls *chartjunk*. Tufte wrote (and researchers, as well as data visualization designers since, have generally agreed) that extraneous elements often hinder, rather than help, our understanding of charts.

Need proof that Scherer and Karamanis's decluttered graph is better than the alternative? Figure ?? shows a version with a few small tweaks to the code to include grid lines and text labels on axes. Prepare yourself for clutter!

Again, it's not just that this cluttered version looks worse. The clutter actively inhibits understanding. Rather than focus on overall drought patterns (the point of the graph), our brain gets stuck reading repetitive and unnecessary axis text.

One of the best ways to reduce clutter is to break a single chart into what are known as *small multiples*. When we look closely at the data visualization, we see that it is not one chart but actually a set of charts. Each rectangle represents one region in one year. If we filter to show the Southwest region in 2003 and add axis titles, we can see in Figure ?? that the x axis shows the week while the y axis shows the percentage of that region at different drought levels.

Zooming in on a single region in a single year also makes the color choices more obvious. The lightest bars show the percentage of the region that is abnormally dry while the darkest bars show the percentage in exceptional drought conditions. These colors, as we'll see shortly, are intentionally chosen to make differences in the drought levels visible to all readers. When I asked Scherer and Karamanis to speak with me about this data visualization, they initially told me that the code for this piece might be too simple to highlight the power of R for data visualization. No, I told them, I want to speak with you precisely because the code is not super complex. The fact that Scherer and Karamanis were able to produce this complex graph with relatively simple code shows the power of R for data visualization. And it is possible *because* of a theory called the grammar of graphics.

The Grammar of Graphics

If you've used Excel to make graphs, you're probably familiar with the menu shown in Figure ?. When working in Excel, your graph-making journey begins by selecting the type of graph you want to make. Want a bar chart? Click the bar chart icon. Want a line chart? Click the line chart icon.

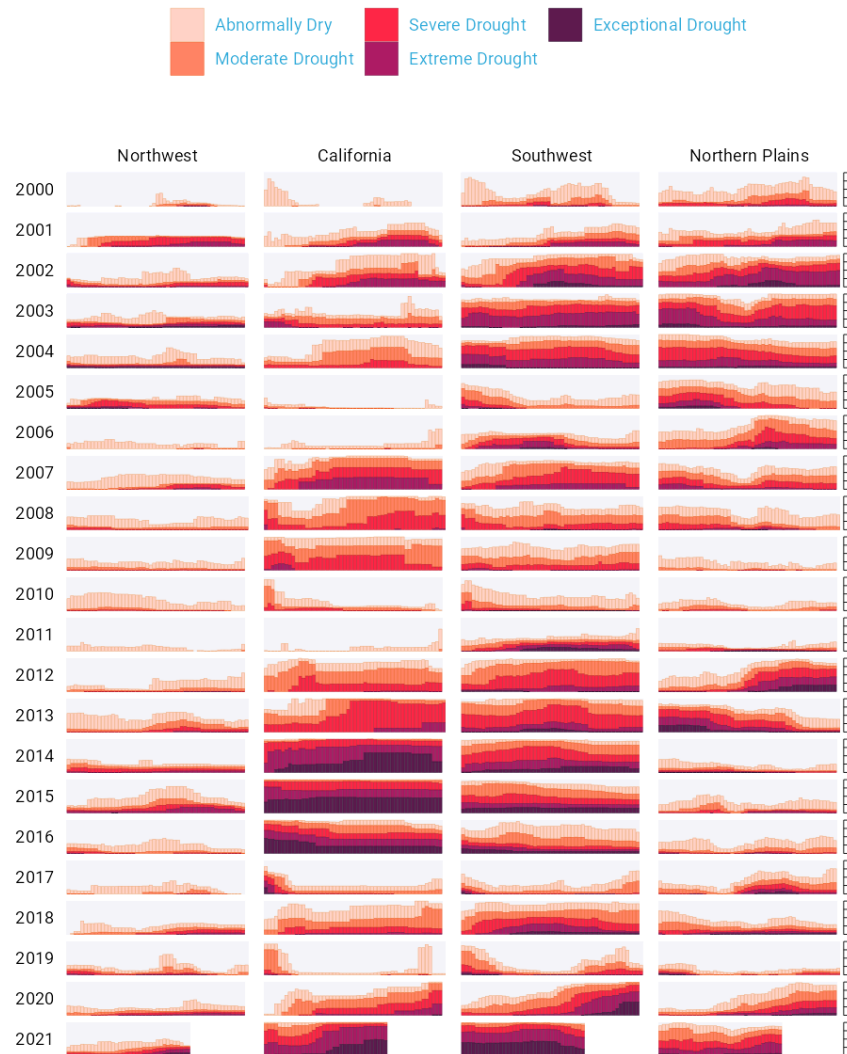


Figure 11: A section of the final drought visualization. If you're incredibly eagle-eyed, you'll see a few minor elements that differ from the version published in *Scientific American*. These are things I had to change to make the plots fit in this book (for example, altering the text size and putting legend text on two rows) or things that *Scientific American* added in post-production (such as annotations).




Figure 12: The cluttered version of the drought visualization




Figure 13: A drought visualization for the Southwest in 2003




Figure 14: The Excel chart chooser menu

f you’ve only ever made data visualization in Excel, this first step may seem so obvious that you’ve never even considered the process of creating data visualization in any other way. But there are different models for thinking about graphs. Rather than conceptualizing graphs types as being distinct, we can recognize the things that they have in common and use these commonalities as the starting point for making them.

This approach to thinking about graphs comes from the late statistician Leland Wilkinson. For years, Wilkinson thought deeply about what data visualization is and how we can describe it. In 1999, he published a book called *The Grammar of Graphics* that sought to develop a consistent way of describing all graphs. In it, Wilkinson argued that we should think of plots not as distinct types à la Excel, but as following a grammar that we can use to describe *any* plot. Just as English grammar tells us that a noun is typically followed by a verb (which is why “he goes” works, while the opposite, “goes he,” does not), knowledge of the grammar of graphics allows us to understand why certain graph types “work.”

Thinking about data visualization through the lens of the grammar of graphics allow us to see, for example, that graphs typically have some data that is plotted on the x axis and other data that is plotted on the y axis. This is the case no matter whether the graph is a bar chart or a line chart, for example. Consider Figure ??, which shows two charts that use identical data on life expectancy in Afghanistan.

While they look different (and would, to the Excel user, be different types of graphs), Wilkinson’s grammar of graphics allows us to see their similarities. (Incidentally, Wilkinson’s feelings on graph-making tools like Excel became clear when he wrote that “most charting packages channel user requests into a rigid array of chart types.”)

When Wilkinson wrote his book, no data visualization tool could implement his grammar of graphics. This would change in 2010, when Hadley Wickham announced the ggplot2 package for R in an article titled “A Layered Grammar of Graphics.” By providing the tools to implement Wilkinson’s ideas, ggplot2 would come to revolutionize the world of data visualization.

Working With ggplot2

The ggplot2 R package (which I, like nearly everyone in the data visualization world, will refer to simply as ggplot) relies on the idea of plots having multiple layers. Let’s walk through some of the most important layers. We’ll begin by selecting variables to map to aesthetic properties. Then we’ll choose a geometric object to use to represent our data. Next we’ll change the aesthetic properties of our chart (the color scheme, for example) using a `scale_` function. And finally we’ll use a `theme_` function to set the overall look-and-feel of our plot.

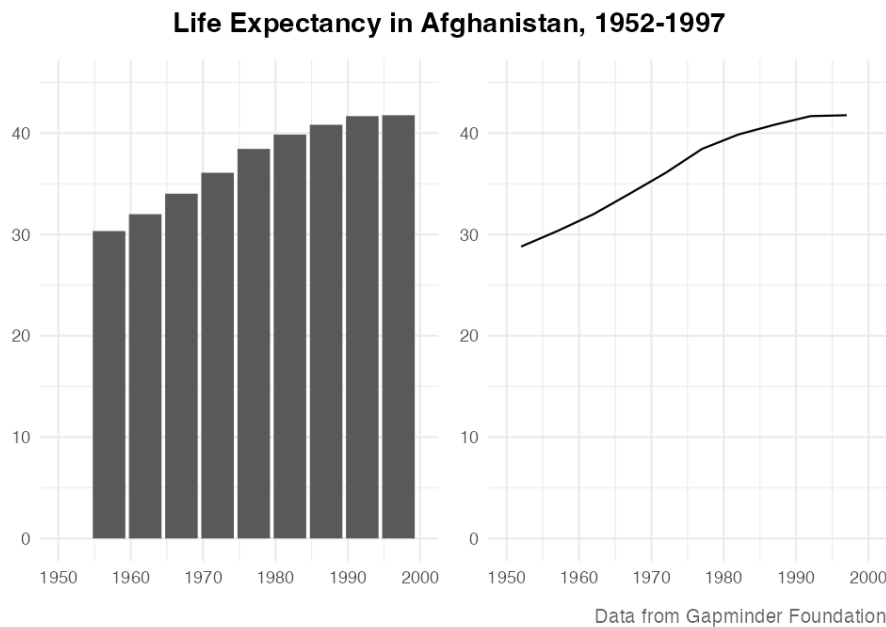


Figure 15: A bar chart and a line chart showing identical data on Afghanistan life expectancy

The First Layer: Mapping Data to Aesthetic Properties

When creating a graph with ggplot, we begin by mapping data to aesthetic properties. All this really means is that we use things like the x or y axis, color, and size (the so-called aesthetic properties) to represent variables. To make this concrete, we'll use the data on life expectancy in Afghanistan, introduced in the previous section, to generate a plot. We can create this data with the following code:

```
library(tidyverse)

gapminder_10_rows <- read_csv("https://data.rwithoutstatistics.com/gapminder_10_rows.csv")
```

Here's what the `gapminder_10_rows` data frame looks like:

```
#> # A tibble: 10 x 6
#>   country    continent  year lifeExp      pop gdpPercap
#>   <chr>      <chr>    <dbl>  <dbl>    <dbl>    <dbl>
#> 1 Afghanistan Asia      1952  28.801  8425333  779.45
#> 2 Afghanistan Asia      1957  30.332  9240934  820.85
```

```
#> 3 Afghanistan Asia      1962  31.997 10267083    853.10
#> 4 Afghanistan Asia      1967  34.02  11537966    836.20
#> 5 Afghanistan Asia      1972  36.088 13079460    739.98
#> 6 Afghanistan Asia      1977  38.438 14880372    786.11
#> 7 Afghanistan Asia      1982  39.854 12881816    978.01
#> 8 Afghanistan Asia      1987  40.822 13867957    852.40
#> 9 Afghanistan Asia      1992  41.674 16317921    649.34
#> 10 Afghanistan Asia     1997  41.763 22227415    635.34
```

If we want to make a chart with `ggplot`, we need to first decide which variable to put on the x axis and which to put on the y axis. Let's say we want to show life expectancy over time. That means we would use the variable `year` on the x axis and the variable `lifeExp` on the y axis. To do so, we begin by using the `ggplot()` function:

```
ggplot(
  data = gapminder_10_rows,
  mapping = aes(
    x = year,
    y = lifeExp
  )
)
```

Within this function, we tell R that we're using the data frame `gapminder_10_rows`. This is the filtered version we created from the full `gapminder` data frame, which includes over 1,700 rows of data. The line following this tells R to use `year` on the x axis and `lifeExp` on the y axis. When we run the code, what we get in Figure ?? doesn't look like much.

But, if you look closely, you can see the beginnings of a plot. Remember that x axis using `year`? There it is! And `lifeExp` on the y axis? Yup, it's there too. I can also see that the values on the x and y axes match up to our data. In the `gapminder_10_rows` data frame, the first year is 1952 and the last year is 1997. The range of the x axis seems to have been created with this data, which goes from 1952 to 1997, in mind (spoiler: it was). And `lifeExp`, which goes from about 28 to about 42 will fit nicely on our y axis.

The Second Layer: Choosing the geoms

Axes are nice, but we're missing any type of visual representation of the data. To get this, we need to add the next layer in `ggplot`: geoms. Short for geometric objects, geoms are functions that provide different ways of representing data. For example, if we want to add points, we use `geom_point()`:




Figure 16: A blank chart

```
ggplot(
  data = gapminder_10_rows,
  mapping = aes(
    x = year,
    y = lifeExp
  )
) +
  geom_point()
```

Now, in Figure ??, we see that people in 1952 had a life expectancy of about 28 and that this value rose through every year in our data.

Let's say we change our mind and want to make a line chart instead. Well, all we have to do is replace `geom_point()` with `geom_line()`:

```
ggplot(
  data = gapminder_10_rows,
  mapping = aes(
    x = year,
    y = lifeExp
  )
)
```




Figure 17: The same chart but with points added

```
) +  
  geom_line()
```

Figure ?? shows the result.

To really get fancy, what if we add both `geom_point()` and `geom_line()`?

```
ggplot(  
  data = gapminder_10_rows,  
  mapping = aes(  
    x = year,  
    y = lifeExp  
  )  
) +  
  geom_point() +  
  geom_line()
```

This code generates a line chart with points, as seen in Figure ??.

We can extend this idea further, as seen in Figure ??, swapping in `geom_col()` to create a bar chart:




Figure 18: The data as a line chart




Figure 19: The data with points and a line

```
ggplot(  
  data = gapminder_10_rows,  
  mapping = aes(  
    x = year,  
    y = lifeExp  
  )  
) +  
  geom_col()
```

Note that the y axis range has been automatically updated, going from 0 to 40 to account for the different geom.




Figure 20: The data as a bar chart

As you can see, the difference between a line chart and a bar chart isn't as great as the Excel chart-type picker might have us think. Both can have the same aesthetic properties (namely, putting years on the x axis and life expectancies on the y axis). They simply use different geometric objects to visually represent the data.

The Third Layer: Altering Aesthetic Properties

Before we return to the drought data visualization, let's look at a few additional layers that can help us alter our bar chart. Say we want to change the color

of our bars. In the grammar of graphics approach to chart-making, this means mapping some variable to the aesthetic property of `fill`. (Slightly confusingly, the aesthetic property of `color` would, for a bar chart, change only the outline of each bar). In the same way that we mapped `year` to the x axis and `y` to `lifeExp`, we can also map `fill` to a variable, such as `year`:

```
ggplot(
  data = gapminder_10_rows,
  mapping = aes(
    x = year,
    y = lifeExp,
    fill = year
  )
) +
  geom_col()
```

The result is shown in Figure ?? . We see now that, for earlier years, the fill is darker, while for later years, it is lighter (the legend, added to the right of our plot, shows this).




Figure 21: The same chart, now with added colors

What if we want to change the fill colors? For that, we use a new *scale layer*. In this case, I'll use the `scale_fill_viridis_c()` function. The `c` at the end

of the function name refers to the fact that the data is continuous, meaning it can take any numeric value:

```
ggplot(
  data = gapminder_10_rows,
  mapping = aes(
    x = year,
    y = lifeExp,
    fill = year
  )
) +
  geom_col() +
  scale_fill_viridis_c()
```

This function changes the default palette to one that is colorblind-friendly and prints well in grayscale. The `scale_fill_viridis_c()` function is just one of many that start with `scale_` and can alter the fill scale.




Figure 22: The same chart with a colorblind-friendly palette

The Fourth Layer: Setting a Theme

A final layer we'll look at is the theme layer. This layer allows us to change the overall look-and-feel of plots (plot backgrounds, grid lines, and so on). Just as